

Blocked algorithms vs. Algorithms by blocks

Paolo Bientinesi
pauldj@cs.umu.se

June 05, 2023
TU Delft




UMEÅ UNIVERSITY

About me

Italy – Tuscany





Computational Geometry



- ▶ Symmetric eigenproblem
 $AX = X\Lambda$

- ▶ FLAME project
Automatic generation of algorithms

Algorithm $LU = A$

Partition $A \rightarrow \left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right), L \rightarrow \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right), U \rightarrow \left(\begin{array}{c|c} U_{TL} & U_{TR} \\ \hline 0 & U_{BR} \end{array} \right)$

where A_{TL}, L_{TL}, U_{TL} , are 0×0

While $m(A_{TL}) \leq m(A)$ **do**

Repartition

$$\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} A_{00} & a_{01} & A_{02} \\ \hline a_{10}^T & \alpha_{11} & a_{12}^T \\ \hline A_{20} & a_{21} & A_{22} \end{array} \right), \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} L_{00} & 0 & 0 \\ \hline l_{10}^T & 1 & 0 \\ \hline L_{20} & l_{21} & L_{22} \end{array} \right),$$

$$\left(\begin{array}{c|c} U_{TL} & U_{TR} \\ \hline 0 & U_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} U_{00} & u_{01} & U_{02} \\ \hline 0 & v_{11} & u_{12}^T \\ \hline 0 & 0 & U_{22} \end{array} \right)$$

where α_{11}, l, v_{11} are scalars

$$v_{11} := \alpha_{11} - l_{10}^T u_{01}$$

$$u_{12}^T := a_{12}^T - l_{10}^T U_{02}$$

$$l_{21} := (a_{21} - L_{20} u_{01}) / v_{11}$$

Continue with


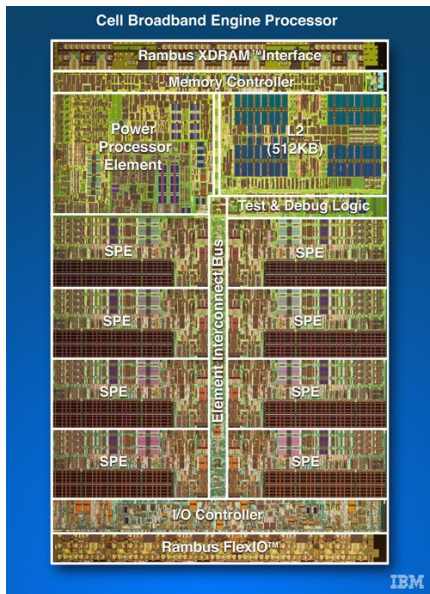
$$\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c|c|c} A_{00} & a_{01} & A_{02} \\ \hline a_{10}^T & \alpha_{11} & a_{12}^T \\ \hline A_{20} & a_{21} & A_{22} \end{array} \right), \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c|c|c} L_{00} & 0 & 0 \\ \hline l_{10}^T & 1 & 0 \\ \hline L_{20} & l_{21} & L_{22} \end{array} \right), \dots$$

endwhile

USA – Duke, NC



- ▶ Cell
- ▶ FFTs



Germany – RWTH Aachen University



High-Performance &
Automatic Computing



Sweden – Umeå University



- ▶ Matrix & tensor operations
- ▶ HPC
- ▶ Computer music



High-Performance
Computing Center
North

Today's outline

- ▶ Part 1: HPC & LA fundamentals
- ▶ Part 2: Unblocked vs. blocked algorithms
- ▶ Part 3: Algorithms by blocks

Fundamentals of HPC & LA

Cholesky: unblocked and blocked algorithms

Cholesky: algorithms by blocks

World of high-performance numerical linear algebra

Where are we?

- ▶ Dense vs. sparse
- ▶ Linear solvers vs. eigensolvers
- ▶ Direct methods vs. iterative methods
- ▶ (Shared memory vs. distributed memory)
- ▶ (Small vs. medium/large size)
- ▶ ...

World of high-performance numerical linear algebra

Where are we?

- ▶ **Dense** vs. sparse
- ▶ **Linear solvers** vs. eigensolvers
- ▶ **Direct methods** vs. iterative methods
- ▶ (**Shared memory** vs. distributed memory)
- ▶ (Small vs. medium/large size)
- ▶ ...

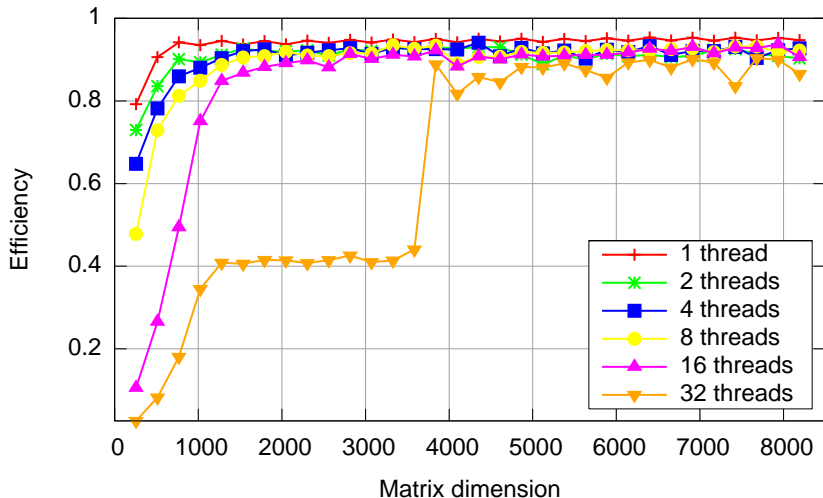
Context: Development of high-performance linear algebra libraries

“Standard” approach (since the '70s):

1. Identify building blocks (“kernels”)
e.g., inner product, matrix-vector & matrix-matrix product, linear system, eigenproblem, ...
2. Encapsulate into a function
e.g., `ddot`, `gemv`, `gemm`, `dgesv`, `dsyevd`, ...
3. Optimize, specialize


Reference: DGEMM – “speed of light”

Efficiency: Percentage of theoretical peak performance



$$\text{GER: } A := A + \alpha xy^T$$

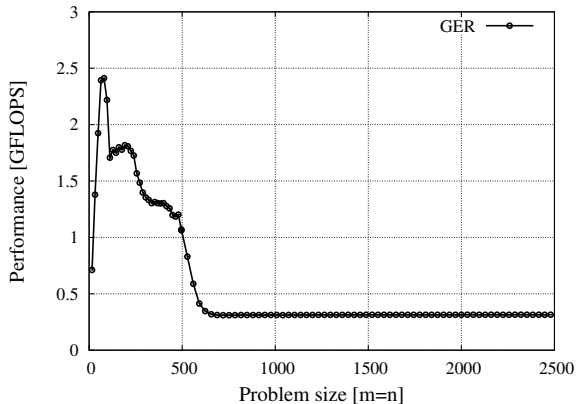
FLOPS = # flops/sec



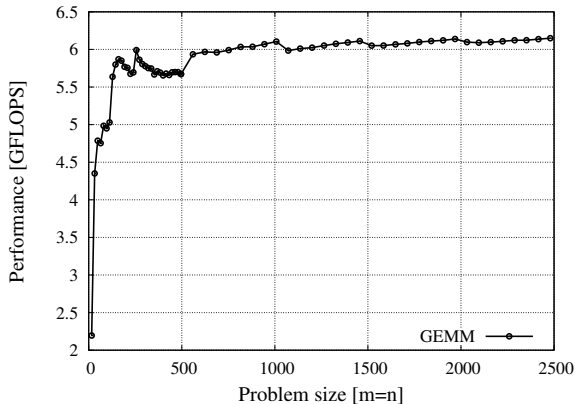
Can you explain these profiles?

$$\text{GER: } A := A + \alpha xy^T$$

FLOPS = # flops/sec



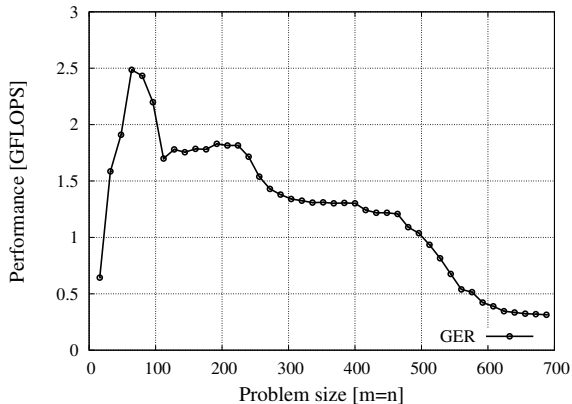
$$\text{GEMM: } C := \alpha A * B + \beta C$$




Can you explain these profiles?

$$\text{GER: } A := A + \alpha xy^T$$

FLOPS = # flops/sec




$$\text{GEMM: } C := \alpha A * B + \beta C$$




Can you explain these profiles?

Memory hierarchy \equiv not all flops cost the same!



Many ways to skin a DGEMM



Can you explain the difference in execution time?

Temporal & spatial locality

Also “Locality of references”, “Principle of locality”

Temporal & spatial locality

Also “Locality of references”, “Principle of locality”

Observations

- ▶ Programs exhibit temporal locality

Consequences

Temporal & spatial locality

Also “Locality of references”, “Principle of locality”

Observations

- ▶ Programs exhibit temporal locality

Consequences

- ▶ Cache memories

Temporal & spatial locality

Also “Locality of references”, “Principle of locality”

Observations

- ▶ Programs exhibit temporal locality
- ▶ Programs exhibit spatial locality

Consequences

- ▶ Cache memories

Temporal & spatial locality

Also “Locality of references”, “Principle of locality”

Observations

- ▶ Programs exhibit temporal locality
- ▶ Programs exhibit spatial locality

Consequences

- ▶ Cache memories
- ▶ Cachelines & Prefetching

Computational intensity

Also “operational intensity”, “arithmetic intensity”

Computational intensity

Also “operational intensity”, “arithmetic intensity”

$$\begin{array}{ll} \text{BLAS-1:} & y := \alpha x + y \qquad x, y \in \mathbb{R}^n \\ & \gamma := \alpha + x^T y \end{array}$$

	#FLOPS	Mem. refs.	Ratio
BLAS-1			

Computational intensity

Also “operational intensity”, “arithmetic intensity”

$$\begin{array}{ll} \text{BLAS-1:} & y := \alpha x + y \\ & \gamma := \alpha + x^T y \end{array} \quad x, y \in \mathbb{R}^n$$

	#FLOPS	Mem. refs.	Ratio
BLAS-1	$2n$		

Computational intensity

Also “operational intensity”, “arithmetic intensity”

$$\begin{array}{ll} \text{BLAS-1:} & y := \alpha x + y \\ & \gamma := \alpha + x^T y \end{array} \quad x, y \in \mathbb{R}^n$$

	#FLOPS	Mem. refs.	Ratio
BLAS-1	$2n$	$3n$	$2/3$

Computational intensity

Also “operational intensity”, “arithmetic intensity”

$$\begin{array}{ll} \text{BLAS-1:} & y := \alpha x + y \quad x, y \in \mathbb{R}^n \\ & \gamma := \alpha + x^T y \end{array}$$

$$\begin{array}{ll} \text{BLAS-2:} & y := Ax + y \quad A, L \in \mathbb{R}^{n \times n}, x, y \in \mathbb{R}^n \\ & y := L^{-1}x \end{array}$$

	#FLOPS	Mem. refs.	Ratio
BLAS-1	$2n$	$3n$	$2/3$
BLAS-2			

Computational intensity

Also “operational intensity”, “arithmetic intensity”

$$\begin{array}{ll} \text{BLAS-1:} & y := \alpha x + y \qquad x, y \in \mathbb{R}^n \\ & \gamma := \alpha + x^T y \end{array}$$

$$\begin{array}{ll} \text{BLAS-2:} & y := Ax + y \quad A, L \in \mathbb{R}^{n \times n}, x, y \in \mathbb{R}^n \\ & y := L^{-1}x \end{array}$$

	#FLOPS	Mem. refs.	Ratio
BLAS-1	$2n$	$3n$	$2/3$
BLAS-2	$2n^2$		

Computational intensity

Also “operational intensity”, “arithmetic intensity”

$$\begin{array}{ll} \text{BLAS-1:} & y := \alpha x + y \qquad x, y \in \mathbb{R}^n \\ & \gamma := \alpha + x^T y \end{array}$$

$$\begin{array}{ll} \text{BLAS-2:} & y := Ax + y \quad A, L \in \mathbb{R}^{n \times n}, x, y \in \mathbb{R}^n \\ & y := L^{-1}x \end{array}$$

	#FLOPS	Mem. refs.	Ratio
BLAS-1	$2n$	$3n$	$2/3$
BLAS-2	$2n^2$	n^2	2

Computational intensity

Also “operational intensity”, “arithmetic intensity”

$$\begin{array}{ll} \text{BLAS-1:} & y := \alpha x + y \quad x, y \in \mathbb{R}^n \\ & \gamma := \alpha + x^T y \end{array}$$

$$\begin{array}{ll} \text{BLAS-2:} & y := Ax + y \quad A, L \in \mathbb{R}^{n \times n}, x, y \in \mathbb{R}^n \\ & y := L^{-1}x \end{array}$$

$$\begin{array}{ll} \text{BLAS-3:} & C := AB + C \quad A, B, C, L \in \mathbb{R}^{n \times n} \\ & C := L^{-1}B \end{array}$$

	#FLOPS	Mem. refs.	Ratio
BLAS-1	$2n$	$3n$	$2/3$
BLAS-2	$2n^2$	n^2	2
BLAS-3			

Computational intensity

Also “operational intensity”, “arithmetic intensity”

BLAS-1:	$y := \alpha x + y$ $\gamma := \alpha + x^T y$	$x, y \in \mathbb{R}^n$
BLAS-2:	$y := Ax + y$ $y := L^{-1}x$	$A, L \in \mathbb{R}^{n \times n}, x, y \in \mathbb{R}^n$
BLAS-3:	$C := AB + C$ $C := L^{-1}B$	$A, B, C, L \in \mathbb{R}^{n \times n}$

	#FLOPS	Mem. refs.	Ratio
BLAS-1	$2n$	$3n$	$2/3$
BLAS-2	$2n^2$	n^2	2
BLAS-3	$2n^3$	$4n^2$	$n/2$

Summary: Libraries

1970s

- ▶ Identification, analysis, optimization of **building blocks**

Summary: Libraries

1970s

- ▶ Identification, analysis, optimization of **building blocks**
- ▶ Computers difficult to program
BUT
Programs “easy” to optimize
- ▶ $\text{Cost}(\mathcal{A}/g) \equiv \# \text{operations}(\mathcal{A}/g)$

Summary: Libraries

1970s

- ▶ Identification, analysis, optimization of **building blocks**
- ▶ Computers difficult to program
BUT
Programs “easy” to optimize
- ▶ $\text{Cost}(\mathcal{A}/g) \equiv \# \text{operations}(\mathcal{A}/g)$
- ▶ Libraries: convenience, portability, separation of concerns, confidence, performance

Summary: Libraries

1970s

- ▶ Identification, analysis, optimization of **building blocks**
- ▶ Computers difficult to program
BUT
Programs “easy” to optimize
- ▶ $\text{Cost}(\mathcal{A}/g) \equiv \# \text{operations}(\mathcal{A}/g)$
- ▶ Libraries: convenience, portability, separation of concerns, confidence, performance

since 1980s

- ▶ Increasingly complex HW:
E.g., vector processors, memory hierarchies, prefetching, ...

Summary: Libraries

1970s

- ▶ Identification, analysis, optimization of **building blocks**
- ▶ Computers difficult to program
BUT
Programs “easy” to optimize
- ▶ $\text{Cost}(\mathcal{A}lg) \equiv \#operations(\mathcal{A}lg)$
- ▶ Libraries: convenience, portability, separation of concerns, confidence, performance

since 1980s

- ▶ Increasingly complex HW:
E.g., vector processors, memory hierarchies, prefetching, ...
- ▶ Computers easier to program
BUT
Programs difficult to optimize
- ▶ $\text{Cost}(\mathcal{A}lg) \neq \#operations(\mathcal{A}lg)$

Summary: Libraries

1970s

- ▶ Identification, analysis, optimization of **building blocks**
- ▶ Computers difficult to program
BUT
Programs “easy” to optimize
- ▶ $\text{Cost}(\mathcal{A}/g) \equiv \# \text{operations}(\mathcal{A}/g)$
- ▶ Libraries: convenience, portability, separation of concerns, confidence, performance

since 1980s

- ▶ Increasingly complex HW:
E.g., vector processors, memory hierarchies, prefetching, ...
- ▶ Computers easier to program
BUT
Programs difficult to optimize
- ▶ $\text{Cost}(\mathcal{A}/g) \neq \# \text{operations}(\mathcal{A}/g)$
- ▶ Libraries: as before +
necessity for performance

Fundamentals of HPC & LA

Cholesky: unblocked and blocked algorithms

Cholesky: algorithms by blocks

Cholesky Factorization

$$LL^T = A \quad L := \Gamma(A)$$

$$L = \left(\begin{array}{c|c} L_{TL} & \\ \hline L_{BL} & L_{BR} \end{array} \right) = ?$$

Cholesky Factorization

$$LL^T = A \quad L := \Gamma(A)$$

$$L = \left(\begin{array}{c|c} L_{TL} & \\ \hline L_{BL} & L_{BR} \end{array} \right) = ?$$

$$\left(\begin{array}{c|c} L_{TL} & \\ \hline L_{BL} & L_{BR} \end{array} \right) \left(\begin{array}{c|c} L_{TL}^T & L_{BL}^T \\ \hline & L_{BR}^T \end{array} \right) = \left(\begin{array}{c|c} A_{TL} & A_{BL}^T \\ \hline A_{BL} & A_{BR} \end{array} \right)$$

Cholesky Factorization

$$LL^T = A \quad L := \Gamma(A)$$

$$L = \left(\begin{array}{c|c} L_{TL} & \\ \hline L_{BL} & L_{BR} \end{array} \right) = ?$$

$$\left(\begin{array}{c|c} L_{TL}L_{TL}^T = A_{TL} & \\ \hline L_{BL}L_{TL}^T = A_{BL} & L_{BL}L_{BL}^T + L_{BR}L_{BR}^T = A_{BR} \end{array} \right)$$

Cholesky Factorization

$$LL^T = A \quad L := \Gamma(A)$$

$$L = \left(\begin{array}{c|c} L_{TL} & \\ \hline L_{BL} & L_{BR} \end{array} \right) = ?$$

Partitioned Matrix Expression (PME):

$$\left(\begin{array}{c|c} L_{TL} = \Gamma(A_{TL}) & \\ \hline L_{BL} = A_{BL} L_{TL}^{-T} & L_{BR} = \Gamma(A_{BR} - L_{BL} L_{BL}^T) \end{array} \right)$$

Cholesky Factorization

$$LL^T = A \quad L := \Gamma(A)$$

$$L = \left(\begin{array}{c|c} L_{TL} & \\ \hline L_{BL} & L_{BR} \end{array} \right) = ?$$

Operations:

$$\left(\begin{array}{c|c} 1) L_{TL} = \text{CHOL} & \\ \hline 2) L_{BL} = \text{TRSM} & 3) L_{BR} = \text{CHOL}(\text{SYRK}) \end{array} \right)$$

Cholesky Factorization

$$LL^T = A \quad L := \Gamma(A)$$


$$L = \left(\begin{array}{c|c} L_{TL} & \\ \hline L_{BL} & L_{BR} \end{array} \right) = ?$$

Dependencies:

$$\left(\begin{array}{c|c} L_{TL} = \Gamma(A_{TL}) & \\ \hline L_{BL} = A_{BL} L_{TL}^{-T} & L_{BR} = \Gamma(A_{BR} - L_{BL} L_{BL}^T) \end{array} \right)$$

Algorithm #1

Iteration i: completed



Algorithm #1

State of the matrix:


$$\left(\begin{array}{c|c} L_{TL} = \text{CHOL} & \end{array} \right)$$

Final state:

$$\left(\begin{array}{c|c} L_{TL} = \text{CHOL} & \\ \hline L_{BL} = \text{TRSM} & L_{BR} = \text{CHOL}(\text{SYRK}) \end{array} \right)$$


Algorithm #1

Iteration i: completed



Algorithm #1

Iteration $i+1$: repartitioning. Blocked vs. unblocked!




Algorithm #1

Iteration $i+1$: computation

trsm	syrk chol	

Algorithm #1


Iteration $i+1$: completed (boundary shift)



A Different Algorithm?

Algorithm #2

Iteration i: completed



Algorithm #2

State of the matrix:


$$\left(\begin{array}{c|c} L_{TL} = \text{CHOL} & \\ \hline L_{BL} = \text{TRSM} & \end{array} \right)$$

Final State:

$$\left(\begin{array}{c|c} L_{TL} = \text{CHOL} & \\ \hline L_{BL} = \text{TRSM} & L_{BR} = \text{CHOL}(\text{SYRK}) \end{array} \right)$$


Algorithm #2

Iteration $i+1$: repartitioning




Algorithm #2

Iteration $i+1$: computation



Algorithm #2

Iteration $i+1$: completed (boundary shift)



Yet Another Algorithm!

Algorithm #3

State of the matrix:


$$\left(\begin{array}{c|c} L_{TL} = \text{CHOL} & \\ \hline L_{BL} = \text{TRSM} & L_{BR} = \text{SYRK} \end{array} \right)$$

Final state:

$$\left(\begin{array}{c|c} L_{TL} = \text{CHOL} & \\ \hline L_{BL} = \text{TRSM} & L_{BR} = \text{CHOL}(\text{SYRK}) \end{array} \right)$$


Algorithm #3

Iteration i: completed




Algorithm #3

Iteration $i+1$: repartitioning




Algorithm #3

Iteration $i+1$: computation




Algorithm #3

Iteration $i+1$: completed (boundary shift)




Algorithm Progression

Iteration i: completed




Algorithm Progression

Iteration $i+1$: repartitioning




Algorithm Progression

Iteration $i+1$: computation




Algorithm Progression

Iteration $i+1$: completed (boundary shift)




Algorithm Progression

Iteration $i+2$: repartitioning




Algorithm Progression

Iteration $i+2$: computation



Algorithm Progression

Iteration $i+2$: complete (boundary shift)



Traditional code

- ▶ C, triple loop, unblocked.

```
for ( j = 0; j < n; j++ )
{
    A[j,j] = sqrt( A[j,j] );

    for ( i = j+1; i < n; i++ )
        A[i,j] = A[i,j] / A[j,j];

    for ( k = j+1; k < n; k++ )
        for ( i = k; i < n; i++ )
            A[i,k] = A[i,k] - A[i,j] * A[k,j];
}
```


Traditional code

► Matlab, blocked.

```
for j = 1:nb:n,
    b = min( n-j+1, nb );

    A(j:j+b-1, j:j+b-1) = Chol( A(j:j+b-1, j:j+b-1) );

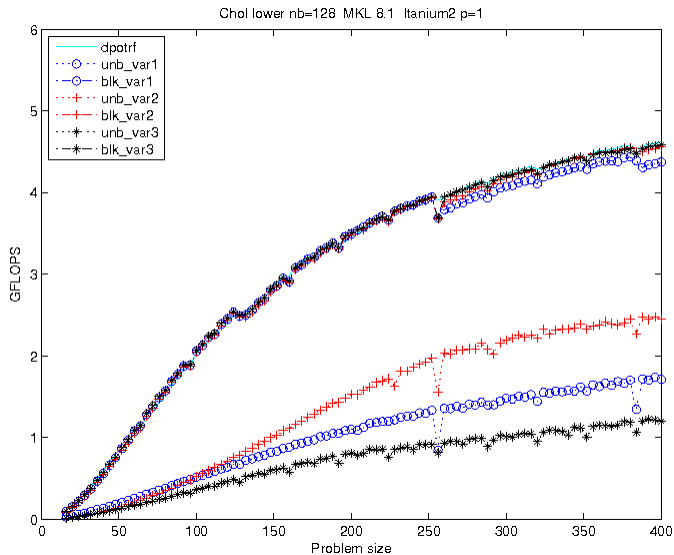
    A(j+b:n, j:j+b-1) = A(j+b:n, j:j+b-1)/A(j:j+b-1, j:j+b-1)';

    A(j+b:n, j+b:n) = A(j+b:n, j+b:n) -
        tril(A(j+b:n, j:j+b-1)) A(j+b:n, j:j+b-1)';
end
```

Traditional code: LAPACK, blocked

```
SUBROUTINE DPOTRF( UPLO, N, A, LDA, INFO )  
[...]  
    DO 20 J = 1, N, NB  
*  
        JB = MIN( NB, N-J+1 )  
        CALL DSYRK( 'Lower', 'No transpose', JB, J-1, -ONE,  
$                A( J, 1 ), LDA, ONE, A( J, J ), LDA )  
        CALL DPOTF2( 'Lower', JB, A( J, J ), LDA, INFO )  
        IF( INFO.NE.0 )  
$            GO TO 30  
        IF( J+JB.LE.N-1 ) THEN  
*  
            CALL DGEMM( 'No transpose', 'Transpose', N-J-JB+1, JB,  
$                J-1, -ONE, A( J+JB, 1 ), LDA, A( J, 1 ),  
$                LDA, ONE, A( J+JB, J ), LDA )  
            CALL DTRSM( 'Right', 'Lower', 'Transpose', 'Non-unit',  
$                N-J-JB+1, JB, ONE, A( J, J ), LDA,  
$                A( J+JB, J ), LDA )  
        END IF
```

Unblocked vs. Blocked Algorithms



Later in Lab #1: LU anyone?

- ▶ Without pivoting:
Precisely the same steps as Cholesky; different PME; different dependencies; 5 algorithms
- ▶ With pivoting:
Same steps, a tad more complicated PME; 3 algorithms


Fundamentals of HPC & LA

Cholesky: unblocked and blocked algorithms


Cholesky: algorithms by blocks

Shared-memory parallelization: Can we do better?

Fork-join \Rightarrow unnecessary synchronizations



Iteration 1



Iteration 2

Synchronization at each iteration; in fact, at each kernel!

Shared-Memory Parallelization

- ▶ Traditional (and pipelined) parallelizations are limited by the dependencies dictated by the code.
- ▶ Parallelism should be limited only by the data dependencies.
- ▶ Idea: imitate a superscalar processor; dynamic detection of data dependencies + out of order execution.

Dependencies (1/3): “True dependency”

Also, “Flow dependency”

$\{x = 1, y = 2, a = 3\}$

...

$y := a * x + y$

$w := 3 * y$

...

$\{y = 5, w = 15\}$

Dependencies (1/3): “True dependency”

Also, “Flow dependency”

$\{x = 1, y = 2, a = 3\}$

...

$y := a * x + y$

$w := 3 * y$

...

$\{y = 5, w = 15\}$

- The value of w depends on the updated value of y

Dependencies (1/3): “True dependency”

Also, “Flow dependency”

$\{x = 1, y = 2, a = 3\}$	
...	...
$y := a * x + y$	$w := 3 * y$
$w := 3 * y$	$y := a * x + y$
...	...
$\{y = 5, w = 15\}$	$\{y = 5, w = 6\}$

- ▶ The value of w depends on the updated value of y
- ▶ The semantics of the program depends on the **order** of the statements

Dependencies (2/3): “Anti dependency”

$\{x = 1, y = 2, a = 3\}$

...

$w := 3 * y$

$y := a * x + y$

...

$\{y = 5, w = 6\}$

- The value of w depends on the initial value of y

Dependencies (2/3): “Anti dependency”

{x = 1, y = 2, a = 3}	
...	...
w := 3 * y	y := a * x + y
y := a * x + y	w := 3 * y
...	...
{y = 5, w = 6}	{y = 5, w = 15}

- ▶ The value of w depends on the initial value of y
- ▶ The semantics of the program depends on the **order** of the statements

Dependencies (3/3): “Output dependency”

Also “Write dependency”

$\{x = 1, y = 2, a = 3\}$

...

$w := 3 * y$

$w := a * x$

...

$\{w = 3\}$

- The value of w depends on the order of the statements

Dependencies (3/3): “Output dependency”

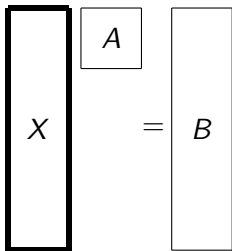
Also “Write dependency”

$\{x = 1, y = 2, a = 3\}$	
...	...
$w := 3 * y$	$w := a * x$
$w := a * x$	$w := 3 * y$
...	...
$\{w = 3\}$	$\{w = 6\}$

- ▶ The value of w depends on the order of the statements
- ▶ The semantics of the program depends on the **order** of the statements

Back to Cholesky: How to create parallelism?


- ▶ Idea: decompose the tasks



The diagram shows a matrix equation $X A = B$. Matrix X is represented by a tall vertical rectangle with a thick black border. Matrix A is a small square located to the right of X . Matrix B is a tall vertical rectangle to the right of A . An equals sign is placed between A and B .

Back to Cholesky: How to create parallelism?

- ▶ Idea: decompose the tasks



Back to Cholesky: How to create parallelism?

- ▶ Idea: decompose the tasks

A diagram illustrating the decomposition of a matrix X . On the left, a tall vertical rectangle labeled X is shown with a thick black border. To its right is an equals sign, followed by a tall vertical rectangle labeled B . Above the B rectangle is a small square labeled A .

A diagram illustrating the decomposition of matrix X into blocks. On the left, a tall vertical rectangle is divided into three horizontal sections labeled X_0 , X_1 , and X_2 from top to bottom. The top section X_0 has a thick black border. To the right of this rectangle is an equals sign, followed by a tall vertical rectangle divided into three horizontal sections labeled B_0 , B_1 , and B_2 from top to bottom. Above the B rectangle is a small square labeled A .

$$X_0 A = B_0$$


$$X_1 A = B_1$$

$$X_2 A = B_2$$


Algorithms by blocks

Also, “Tiled algorithms”. Not “blocked”!

Goal: Create small tasks, feed all processors as early as possible



Iteration i



Iteration i+1

Breaking down the computation

Decomposition in tiles (iteration 1)

chol			
trsm	syrk		
trsm	gemm	syrk	
trsm	gemm	gemm	syrk


Breaking down the computation

Decomposition in tiles (iteration 2)

	chol		
	trsm	syrk	
	trsm	gemm	syrk


Breaking down the computation

Decomposition in tiles (iteration 3)



Breaking down the computation

Decomposition in tiles (iteration 4)



Dependencies

chol			
trsm	syrk		
trsm	gemm	syrk	
trsm	gemm	gemm	syrk

Dependencies

chol			
trsm	syrk		
trsm	gemm	syrk	
trsm	gemm	gemm	syrk

Dependencies

chol			
trsm	syrk		
trsm	gemm	syrk	
trsm	gemm	gemm	syrk

Dependencies

chol			
trsm	syrk		
trsm	gemm	syrk	
trsm	gemm	gemm	syrk

Dependencies

chol			
trsm	syrk		
trsm	gemm	syrk	
trsm	gemm	gemm	syrk

Dependencies

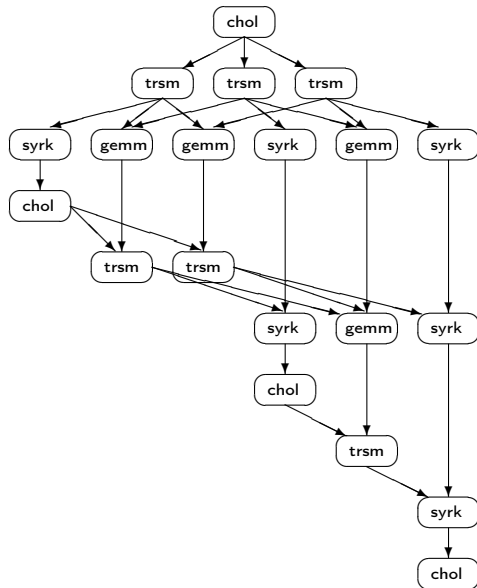
chol			
trsm	syrk		
trsm	gemm	syrk	
trsm	gemm	gemm	syrk

Dependencies


chol			
trsm	syrk		
trsm	gemm	syrk	
trsm	gemm	gemm	syrk

DAG - Dependencies

4×4 -tile matrix



Crossover, 16 cores



Runtime systems, e.g., SuperMatrix, StarPU, Quark, ...

Taskqueues

- ▶ The runtime system “pre-executes” the code.
Whenever a kernel is encountered, one or more tasks are created and inserted in a global task queue.

Runtime systems, e.g., SuperMatrix, StarPU, Quark, ...

Taskqueues

- ▶ The runtime system “pre-executes” the code.
Whenever a kernel is encountered, one or more tasks are created and inserted in a global task queue.
- ▶ Dependencies between tasks are calculated, forming a DAG.

Runtime systems, e.g., SuperMatrix, StarPU, Quark, ...

Taskqueues

- ▶ The runtime system “pre-executes” the code.
Whenever a kernel is encountered, one or more tasks are created and inserted in a global task queue.
- ▶ Dependencies between tasks are calculated, forming a DAG.
- ▶ Tasks with all input operands available are executable.
The other tasks must wait in the queue.

Runtime systems, e.g., SuperMatrix, StarPU, Quark, ...

Taskqueues

- ▶ The runtime system “pre-executes” the code.
Whenever a kernel is encountered, one or more tasks are created and inserted in a global task queue.
- ▶ Dependencies between tasks are calculated, forming a DAG.
- ▶ Tasks with all input operands available are executable.
The other tasks must wait in the queue.
- ▶ Threads asynchronously dequeue tasks from the queue.

Runtime systems, e.g., SuperMatrix, StarPU, Quark, ...

Taskqueues

- ▶ The runtime system “pre-executes” the code.
Whenever a kernel is encountered, one or more tasks are created and inserted in a global task queue.
- ▶ Dependencies between tasks are calculated, forming a DAG.
- ▶ Tasks with all input operands available are executable.
The other tasks must wait in the queue.
- ▶ Threads asynchronously dequeue tasks from the queue.
- ▶ Upon termination of a task, the thread notifies dependent tasks and updates the queue.

Runtime systems, e.g., SuperMatrix, StarPU, Quark, ...

Taskqueues

- ▶ The runtime system “pre-executes” the code.
Whenever a kernel is encountered, one or more tasks are created and inserted in a global task queue.
- ▶ Dependencies between tasks are calculated, forming a DAG.
- ▶ Tasks with all input operands available are executable.
The other tasks must wait in the queue.
- ▶ Threads asynchronously dequeue tasks from the queue.
- ▶ Upon termination of a task, the thread notifies dependent tasks and updates the queue.
- ▶ Loop until all tasks complete execution.

Task Execution

4 × 4-tile matrix

Stage	Scheduled Tasks			
1	chol			
2	trsm	trsm	trsm	trsm
3	syrk	gemm	syrk	gemm
4	gemm	syrk	gemm	gemm
5	gemm	syrk	chol	
6	trsm	trsm	trsm	
7	syrk	gemm	syrk	gemm
8	gemm	syrk	chol	
9	trsm	trsm		
10	syrk	gemm	syrk	
11	chol			
12	trsm			
13	syrk			
14	chol			

SPD Inv: 1) Chol 2) Inv 3) Mat Mat Mult.

5×5 -tile matrix

$$A := A^{-1}$$

$$A := (LL^T)^{-1}$$

$$A := L^{-T}L^{-1}$$

SPD Inv: 1) Chol 2) Inv 3) Mat Mat Mult.

5 × 5-tile matrix

Stage	Scheduled Tasks			
1	chol			
2	trsm	trsm	trsm	trsm
3	syrk	gemm	syrk	gemm
4	gemm	syrk	gemm	gemm
5	gemm	syrk	chol	trsm
6	trsm	trsm	trsm	trsm
7	trsm	trsm	trinv	syrk
8	gemm	syrk	gemm	gemm
9	syrk	ttmm	chol	trsm
10	trsm	trsm	trsm	trsm
11	gemm	gemm	gemm	syrk
12	gemm	syrk	trsm	chol
13	trsm	trsm	trinv	syrk
14	trsm	gemm	gemm	gemm
15	gemm	trmm	syrk	trsm
16	trsm	ttmm	chol	trsm
17	syrk	trinv	gemm	syrk
18	gemm	gemm	gemm	trmm
19	trmm	trsm	trsm	trsm
20	trsm	trsm	trsm	trsm
21	ttmm	syrk	gemm	syrk
22	trinv	gemm	gemm	trinv
23	syrk	syrk	gemm	syrk
24	trmm	gemm	trmm	gemm
25	trmm	syrk	gemm	gemm
26	ttmm	gemm	trmm	trmm
27	syrk	trmm		
28	trmm			
29	ttmm			

Not quite so simple

Not quite so simple

- ▶ Storage by blocks

Not quite so simple

- ▶ Storage by blocks
- ▶ Critical path


Not quite so simple

- ▶ Storage by blocks
- ▶ Critical path
- ▶ Cache “simulator”

Not quite so simple

- ▶ Storage by blocks
- ▶ Critical path
- ▶ Cache “simulator”
- ▶ Tension between size of blocks and number of blocks

SPD Inverse, 16 cores



Multithreaded BLAS vs. Algorithms-by-blocks

No absolute winner: crossover!

✓ Ease of use

✗ Synchronization

✓ Out of order execution

✓ Parallelism dictated by data dependencies

✗ Plateaux